

ETC3250/5250: Introduction to Machine Learning

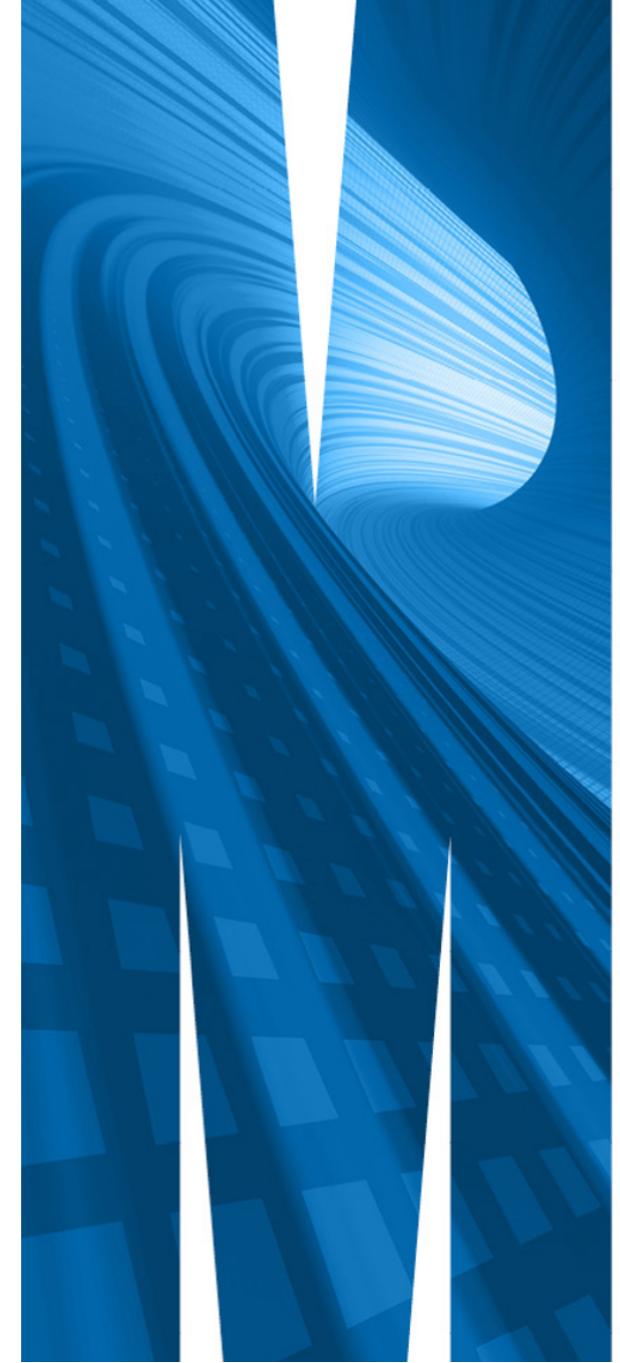
Conceptual framework

Lecturer: Professor Di Cook

Department of Econometrics and Business Statistics

✉ ETC3250.Clayton-x@monash.edu

CALENDAR
Week 1b



Two purposes

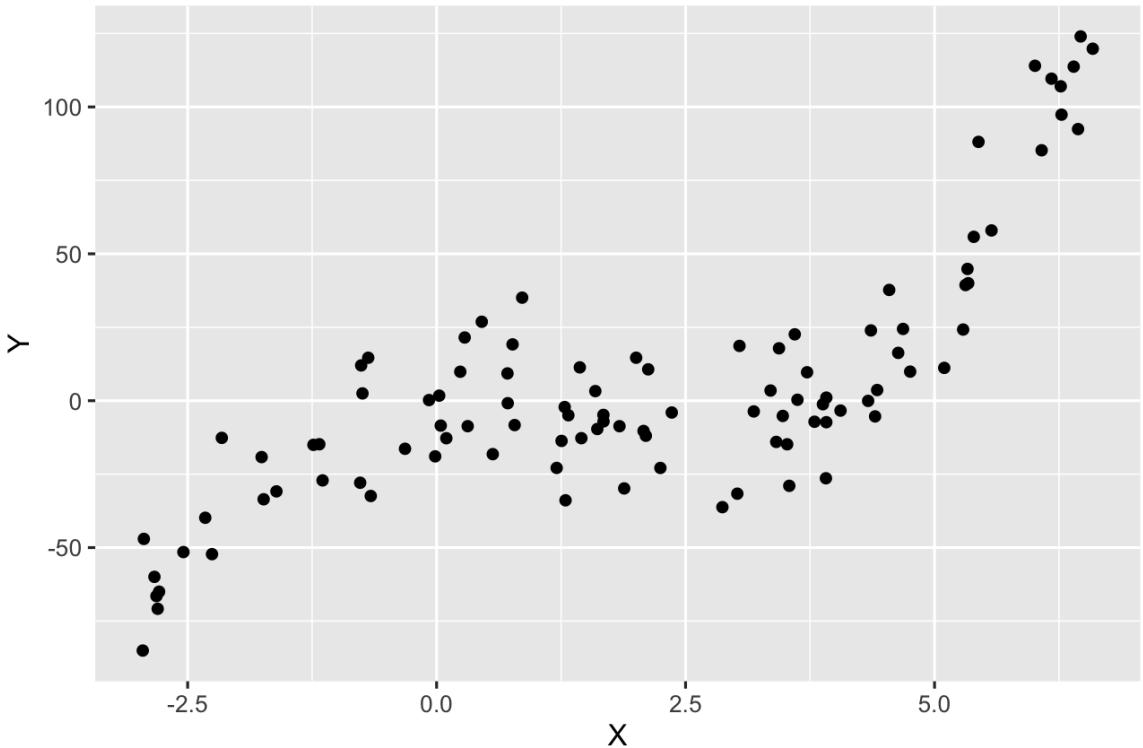
Prediction and inference (or understanding)

Supervised learning

We assume that there is a relationship between Y and \mathbf{X} that can be written as

$$Y = f(\mathbf{X}) + \varepsilon$$

where function $f(\mathbf{X})$ is **fixed but unknown**, and ε is independent of \mathbf{X} and has mean 0.



Note: Data is simulated from a known f .

Why estimate f ?

In many situations, \mathbf{X} may be readily available, but Y ...
might be hard to collect. SO, we would like to be able
to use \mathbf{X} to **predict** new values of Y .

We might not so much be concerned about whether f
is easy to understand, just that we are confident that
its going to do a good job of predicting new values.

$$\hat{Y} = \hat{f}(\mathbf{X})$$

where \wedge reflects what we estimate, the unknown
function to be, and the estimated response value.

Why estimate f ?

In many situations, \mathbf{X} may be readily available, but Y might be hard to collect. So, we would like to be able to use \mathbf{X} to predict new values of Y .

We might not so much be concerned about whether f is easy to understand, just that we are confident that it's going to do a good job of predicting new values.

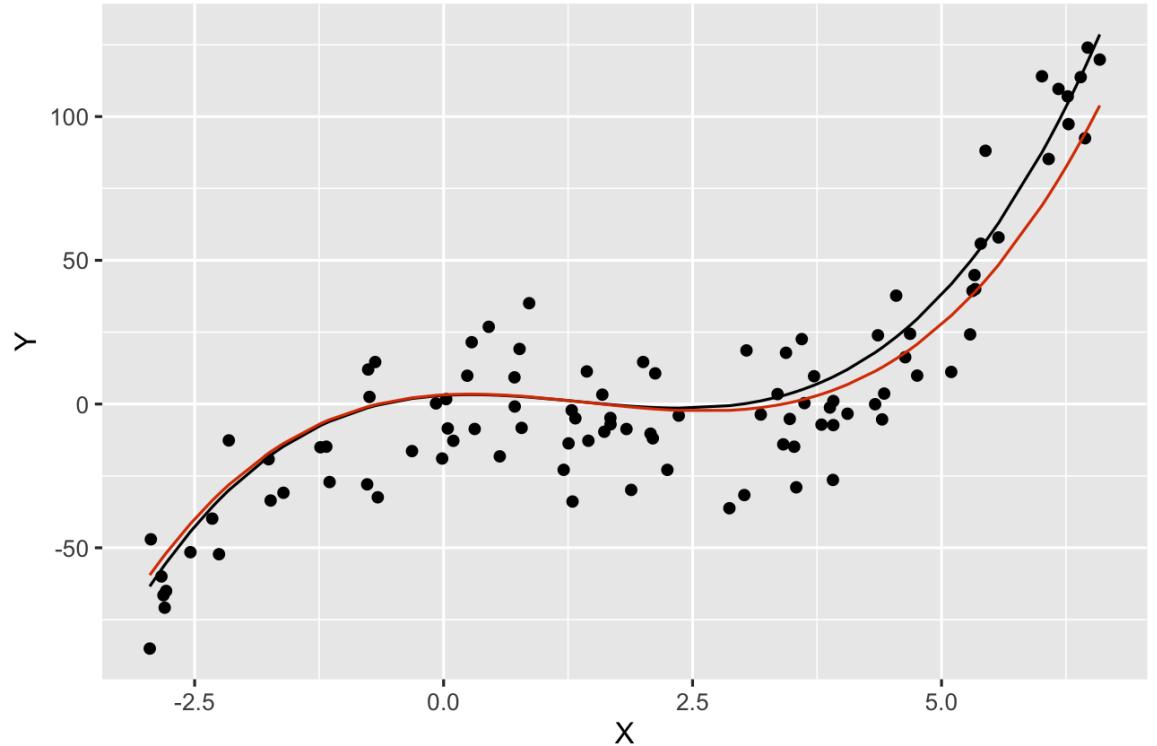
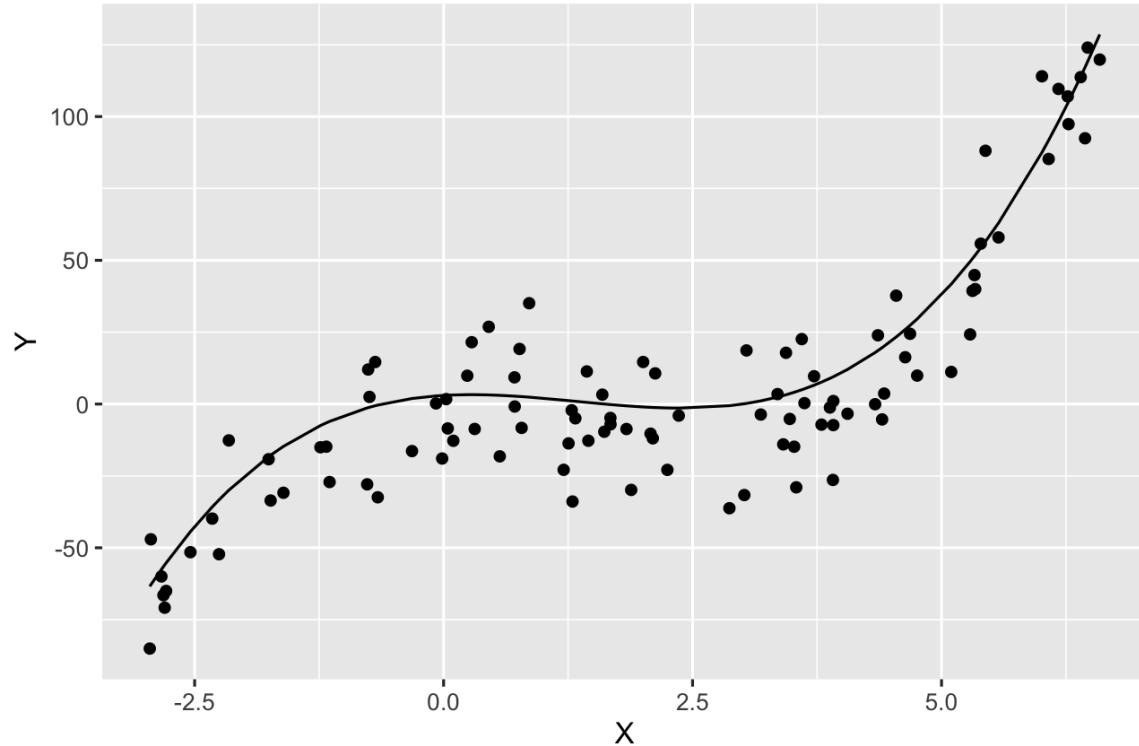
$$\hat{Y} = \hat{f}(\mathbf{X})$$

where $\hat{\cdot}$ reflects what we estimate, the unknown function to be, and the estimated response value.

The accuracy of \hat{Y} as a prediction for Y depends on what we will call reducible AND irreducible error. We can write this as

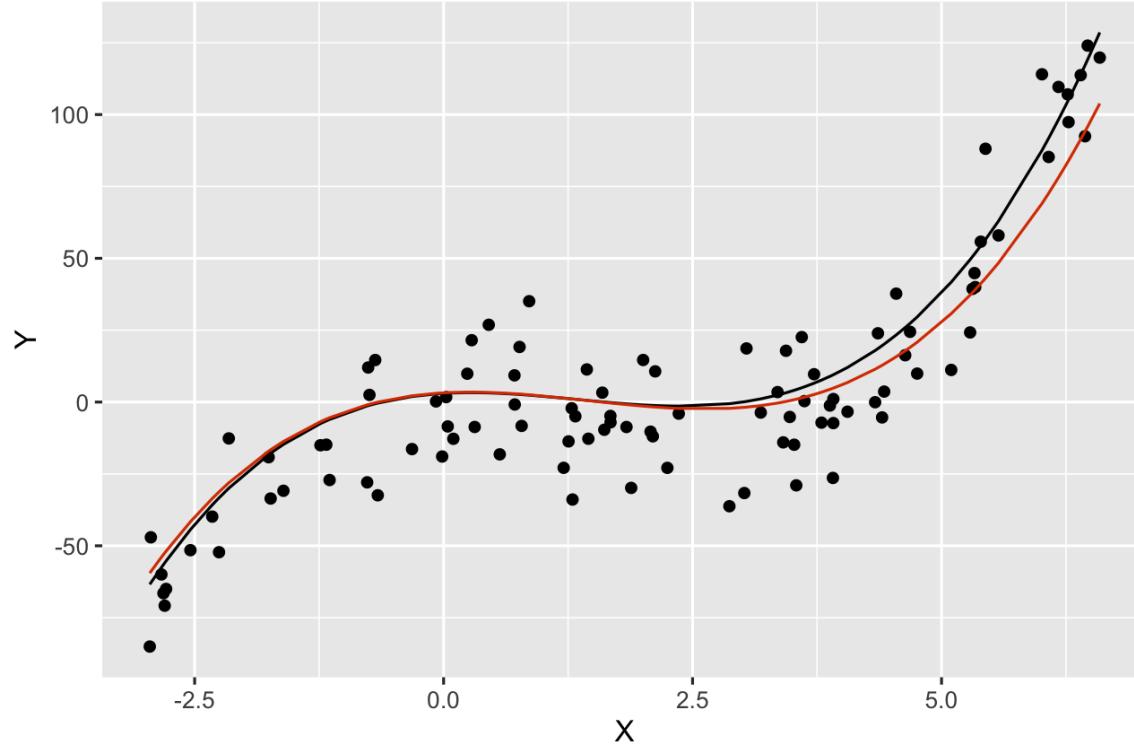
$$\begin{aligned} E(Y - \hat{Y})^2 &= E(f(\mathbf{X}) + \varepsilon - \hat{f}(\mathbf{X}))^2 \\ &= \underbrace{E(f(\mathbf{X}) - \hat{f}(\mathbf{X}))^2}_{\text{reducible}} + \underbrace{\text{Var}(\varepsilon)}_{\text{irreducible}} \end{aligned}$$

where $E(Y - \hat{Y})^2$ represents the average or expected value of the squared difference between the observed and predicted response, and $\text{Var}(\varepsilon)$ represents the variance of the error.



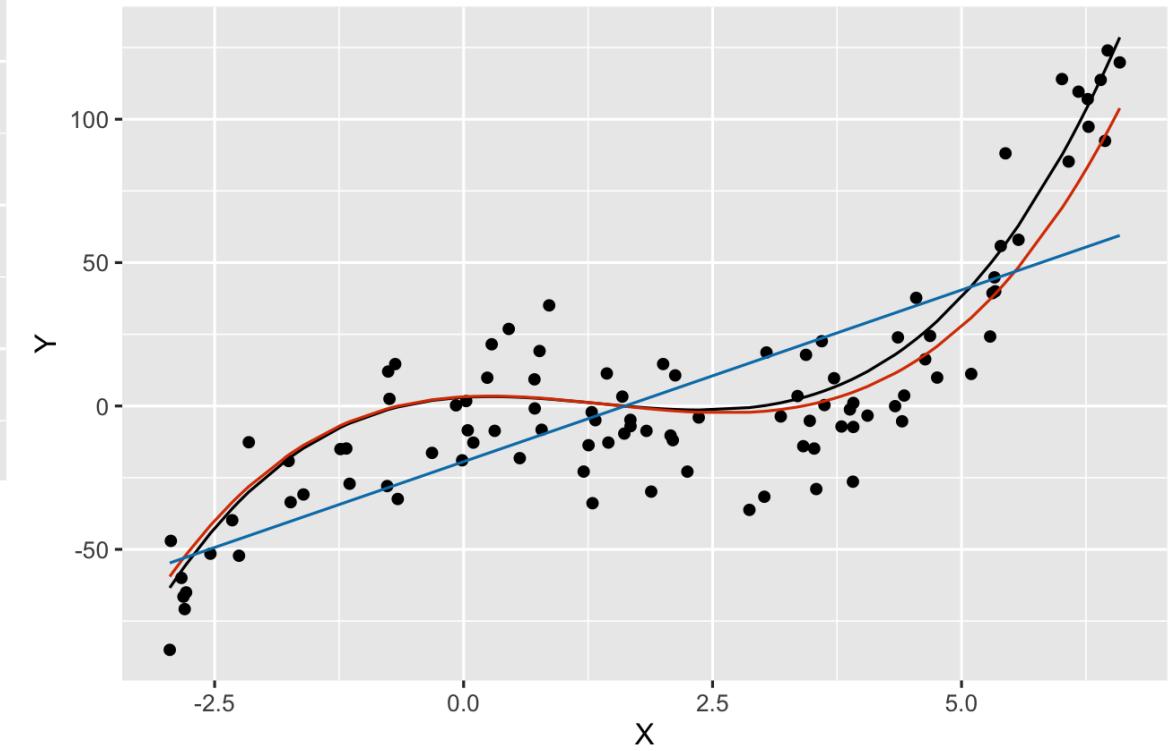
Line indicates **true f** . This is the best possible because **Orange** line indicates an **estimated model \hat{f}** . This fit is all that is left is irreducible error.

very similar to true f but has some room for improvement.



Orange line indicates an estimated model \hat{f} . This fit is very similar to true f but has some room for improvement.

Suppose we used a much simpler model, a linear model.



Blue line indicates a simpler estimated model \hat{f} . There is a lot of room to improve this.

Remember



reducible is what we can **improve** on by producing the **best model**.



irreducible there is some random fluctuation from one sample to the next which is not systematic.



The goal is that the predictions from the model are accurate for future samples.

Inference (understanding)

Inference

We would like to understand the way that Y is related to \mathbf{X} .

- Which predictors are associated with the response?
- What is the relationship between the response and each predictor?
- Can the relationship between Y and each predictor be adequately summarized using a linear equation, or is the relationship more complicated?

It is important here not to treat f as a black box.



Ideally, good prediction also allows for good inference and understanding.

How do we estimate f ?

- Parametric methods:

- Assume that the model takes a specific form
- Fitting then is a matter of estimating the parameters of the model
- Generally considered to be less flexible
- If assumptions are wrong, performance likely to be poor

- Non-parametric methods:

- No specific assumptions
- Allow the data to specify the model form, without being too rough or wiggly
- More flexible
- Generally needs more observations

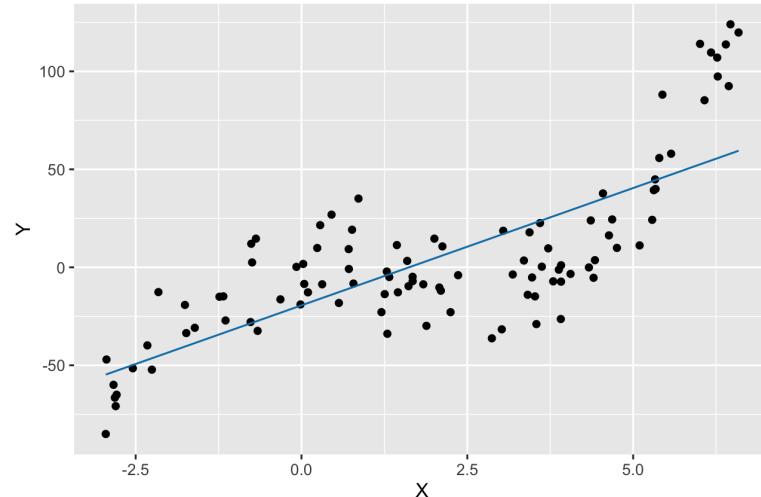
Parametric models

When you force your data to fit the constraints of your model



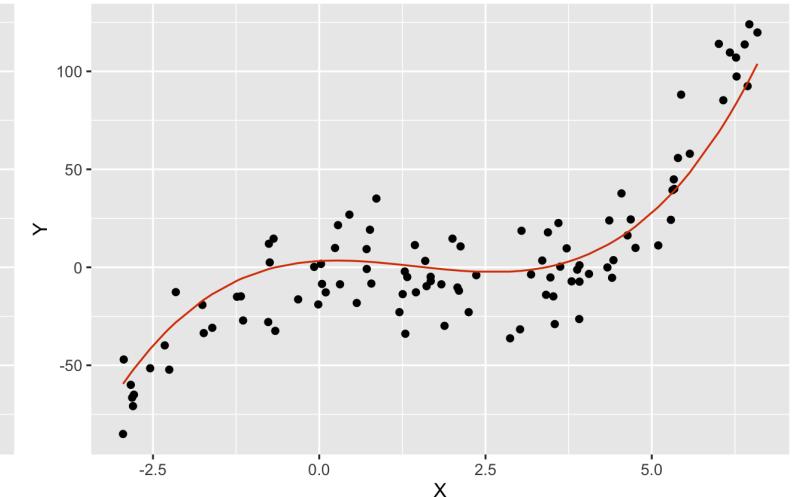
A cat that looks like a bowl.

A linear regression model:
 $f(\mathbf{X}) = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p$



Less flexible

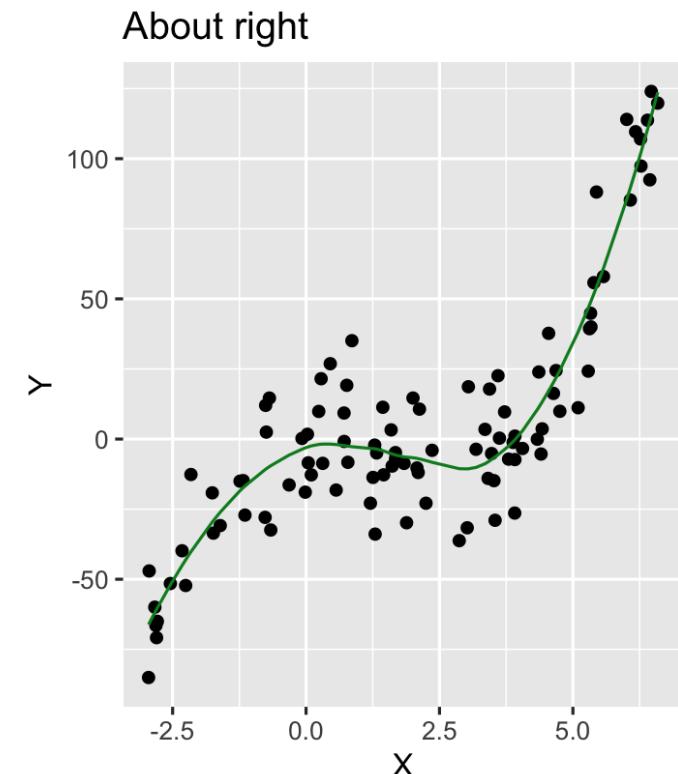
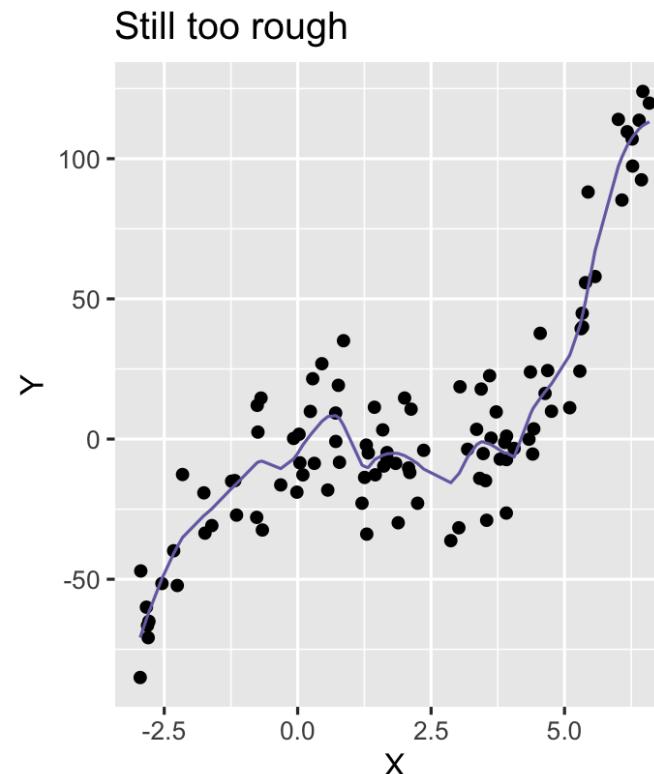
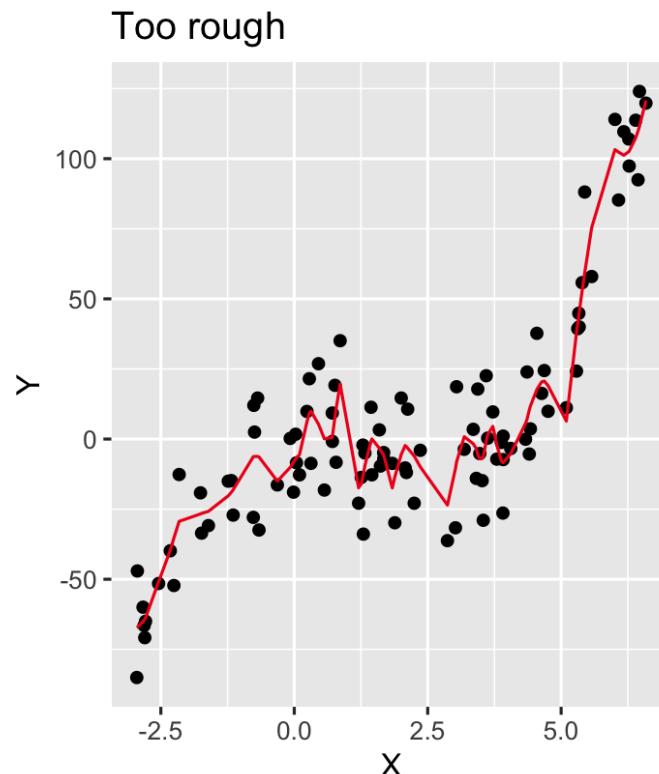
And nonlinear regression model:
 $f(\mathbf{X}) = \beta_0 + \beta_1 X_1 + \beta_2 X_1^2 + \dots$



More flexible

Non-parametric models (even more flexible)

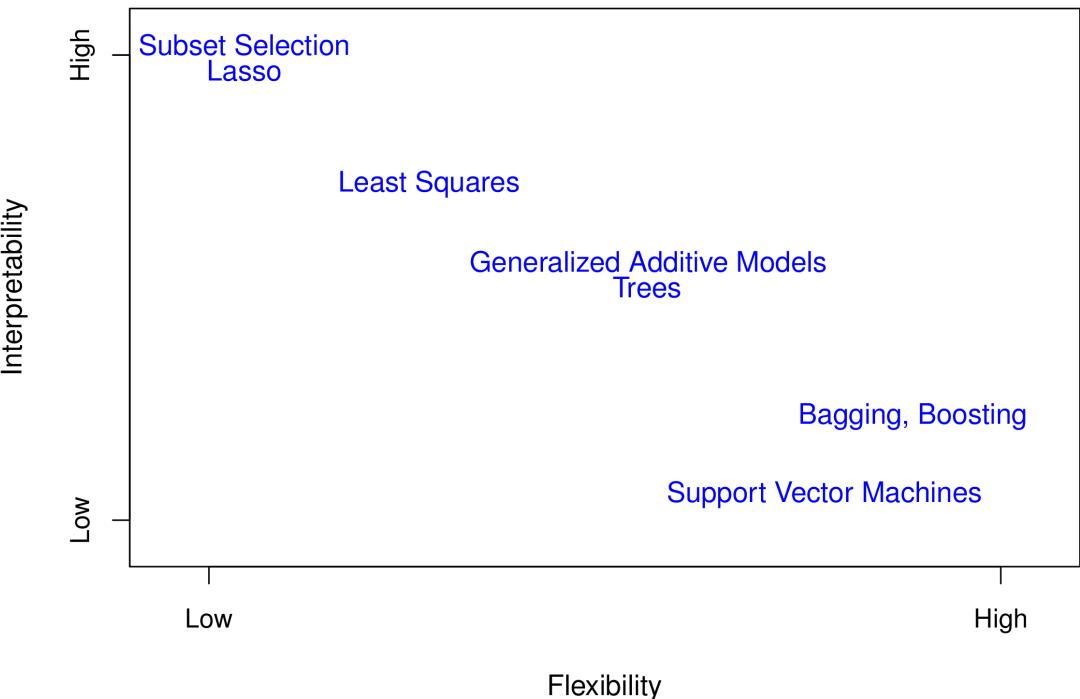
Example: Local polynomial regression, called loess. Fit a linear model to many small subsets of the data.



A more general approach is called k -nearest neighbours, $\hat{Y}(x) = \frac{1}{k} \sum_{x_i \in n_k(x)} y_i$.

Trade-off between predictive accuracy and model interpretability

A summary of common models and how they tend to lie in terms of predictive accuracy vs interpretability.



Assessing model accuracy

Assessing model accuracy



Training data: the set of observations used to train or teach the method to estimate f .

Suppose we have a regression model $y = f(\mathbf{x}) + \varepsilon$. Estimate \hat{f} from some **training data**, $Tr = \{(y_i, \mathbf{x}_i)\}_{i=1}^n$.

The most common measure of accuracy is the **training Mean Squared Error (MSE)**

$$MSE_{Tr} = \text{Ave}_{i \in Tr} [y_i - \hat{f}(\mathbf{x}_i)]^2 = \frac{1}{n} \sum_{i=1}^n [(y_i - \hat{f}(\mathbf{x}_i))^2]$$

Assessing model accuracy



Test data: the set of observations reserved to compute accuracy of the model for **new** data.

A better measure of **accuracy** is obtained by using the **test data**, denoted as $Te = \{(y_i, \mathbf{x}_i)\}_{i=1}^m$, **Test Mean Squared Error**

$$MSE_{Te} = \text{Ave}_{j \in Te} [y_j - \hat{f}(\mathbf{x}_j)]^2 = \frac{1}{m} \sum_{j=1}^m [(y_j - \hat{f}(\mathbf{x}_j))^2]$$

Regression vs classification



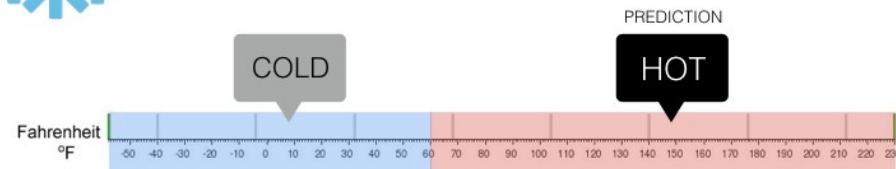
Regression

What is the temperature going to be tomorrow?



Classification

Will it be Cold or Hot tomorrow?



Source: Taylor Fogarty

Assessing model accuracy for classification



To indicate the categorical response, we will use \hat{C} instead of \hat{f} .

Compute \hat{C} from some **training data**, $Tr = \{(y_i, \mathbf{x}_i)\}_{i=1}^n$. In place of MSE, we now use the error rate (**fraction of misclassifications**) to get the **Training Error Rate**

$$\text{Error rate}_{Tr} = \frac{1}{n} \sum_{i=1}^n I(y_i \neq \hat{C}(\mathbf{x}_i))$$

And again a better estimate of future **accuracy** is obtained using **test data** $Te = \{(y_i, \mathbf{x}_i)\}_{i=1}^m$ to get the **Test Error Rate**

$$\text{Error rate}_{Te} = \frac{1}{m} \sum_{j=1}^m I(y_j \neq \hat{C}(\mathbf{x}_j))$$



Generally, training error will be smaller than test error.

Because the training data is used to fit the model, by design the error will be small relative to the error when the model is used on new data.

Bias-variance trade-off



There are two competing forces that govern the choice of learning method: **bias** and **variance**.

Bias is the error that is introduced by modeling a complicated problem by a simpler problem.

- For example, linear regression assumes a linear relationship and perhaps the relationship is not exactly linear.
- In general, but not always, the **more flexible** a method is, the **less bias** it will have.

[This site](#) has a lovely explanation.

Bias-variance trade-off



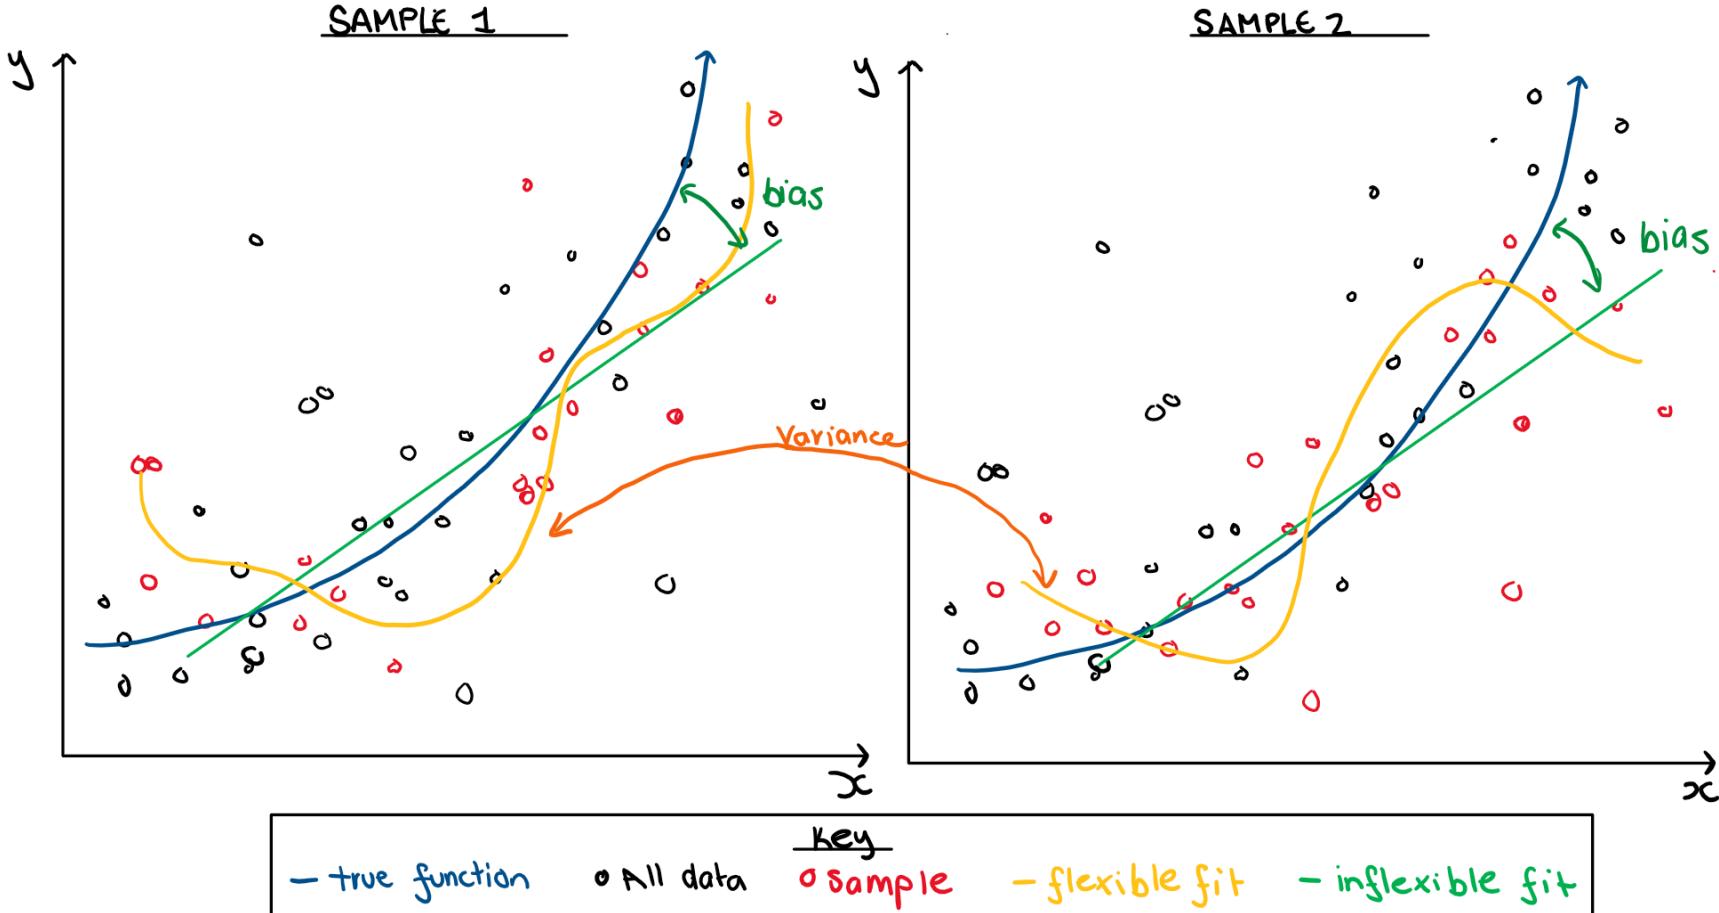
There are two competing forces that govern the choice of learning method: **bias** and **variance**.

Variance refers to how much your estimate would change if you had different training data. Its measuring how much your model depends on the data you have, to the neglect of future data.

- In general, the **more flexible** a method is, the **more variance** it has.
- The **size** of the training data has an impact on the variance.

Flexibility, bias and variance

This blog post by Harriet Mason, former ETC3250 student has a lovely explanation of the trade-off in flexibility and effect on bias and variance.



MSE decomposition into bias and variance

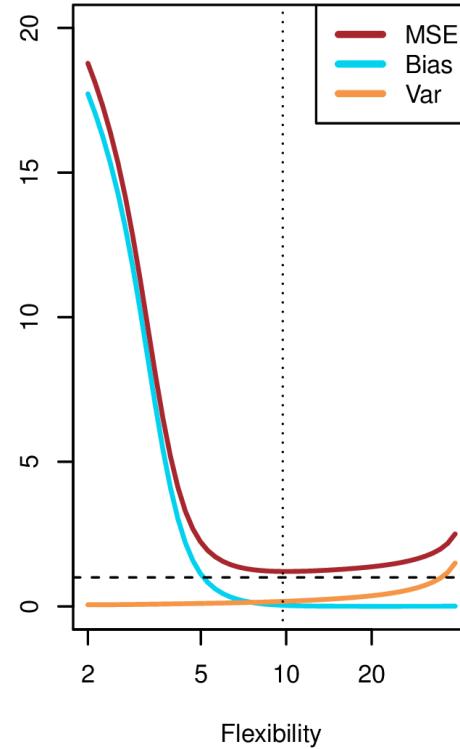
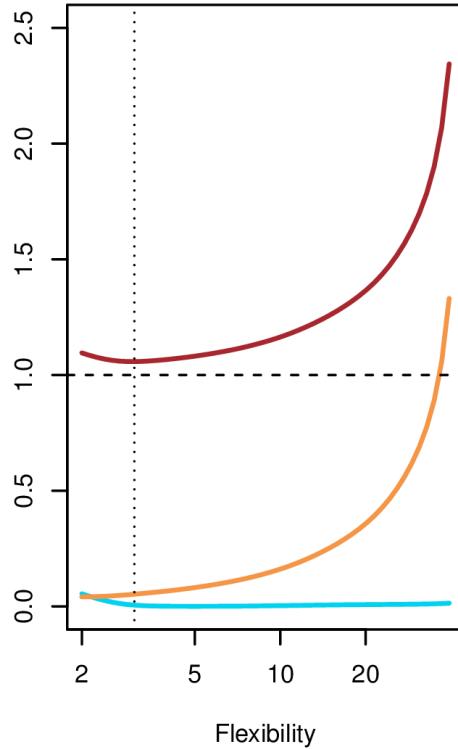
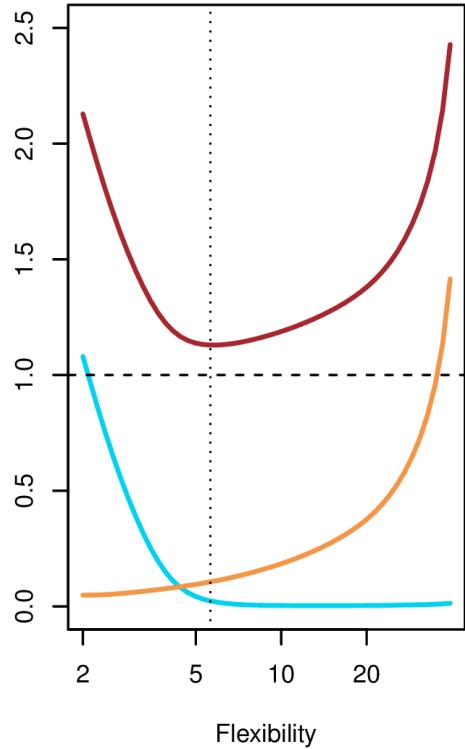
The expected **test** MSE for a new y_0 at a new observation, called \mathbf{x}_0 , will be equal to

$$E[(y_0 - \hat{f}(\mathbf{x}_0))^2] = [\text{Bias}(\hat{f}(\mathbf{x}_0))]^2 + \text{Var}(\hat{f}(\mathbf{x}_0)) + \text{Var}(\varepsilon)$$

Test MSE = Bias² + Variance + Irreducible variance

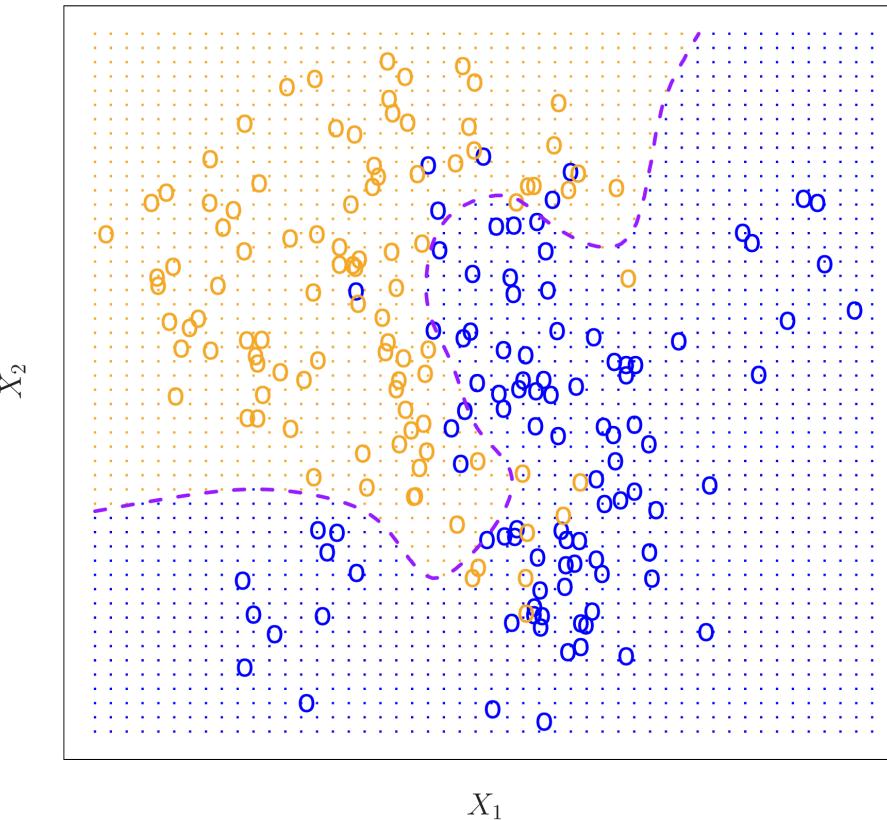
- The expectation averages over the variability of Y as well as the variability in the training data.
- As the flexibility of \hat{f} increases, its variance increases and its bias decreases.
- To decide on the best model, from a range with different flexibility, you choose based on average test MSE at the **bias-variance trade-off**, where both are minimised.

Bias-variance tradeoff



Squared bias, variance, $\text{Var}(\varepsilon)$ (dashed line), and test MSE for the three data sets shown earlier. The vertical dotted line indicates the flexibility level corresponding to the smallest test MSE.

A case study using nearest neighbours classification



Colour indicates true class of each observation.
Dashed line indicates true boundary.

K Nearest Neighbours (KNN)

One of the simplest classifiers. Given a test observation \mathbf{x}_0 ,

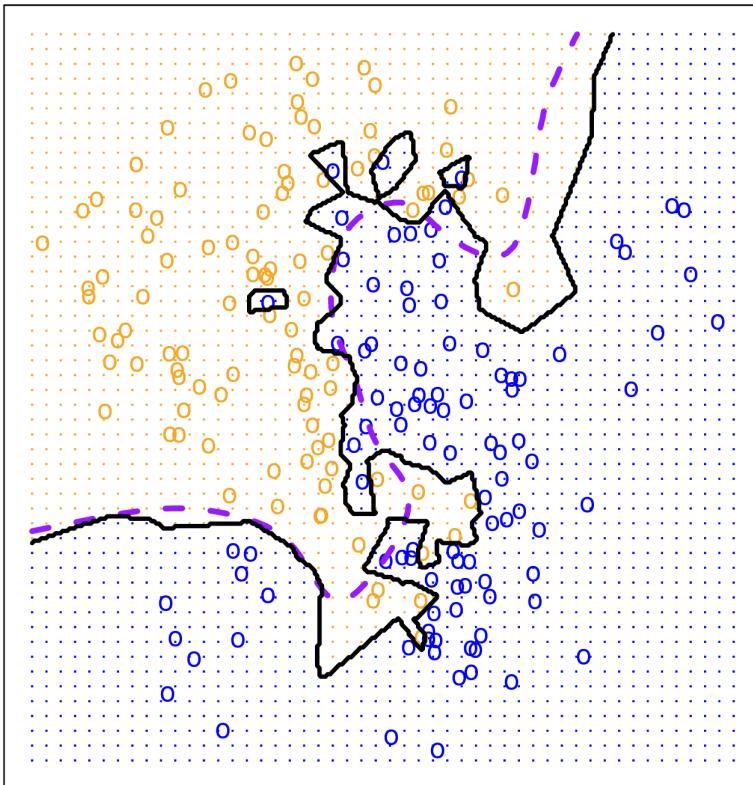
- Find the K nearest points to \mathbf{x}_0 in the training data, call this \mathcal{N}_0 .
- Estimate conditional probabilities

$$P(Y = C_j \mid \mathbf{X} = \mathbf{x}_0) = \frac{1}{K} \sum_{i \in \mathcal{N}_0} I(y_i = C_j).$$

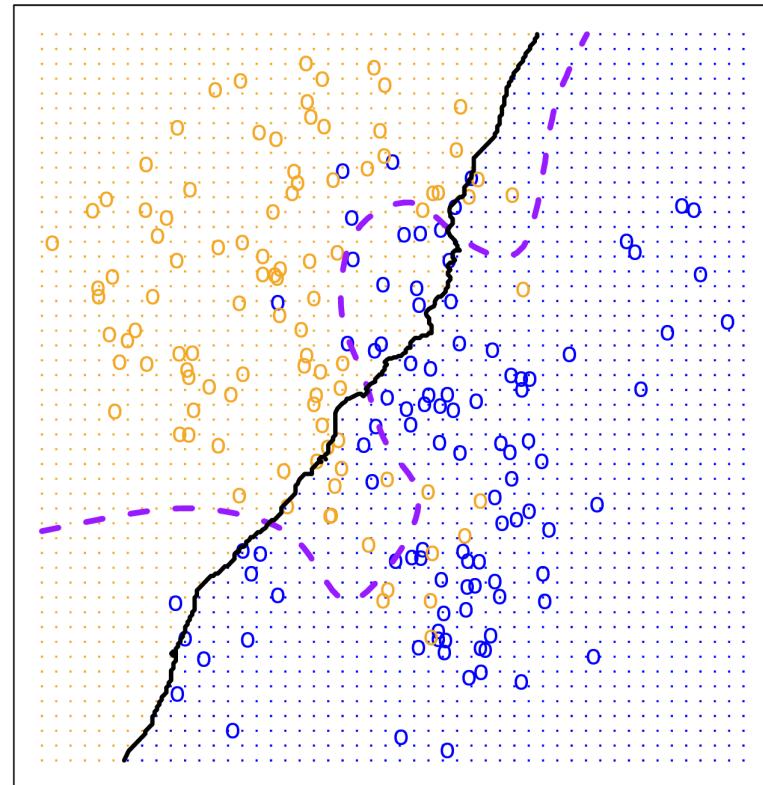
- Classify \mathbf{x}_0 to class with largest probability.

KNN: too flexible and not enough

KNN: K=1



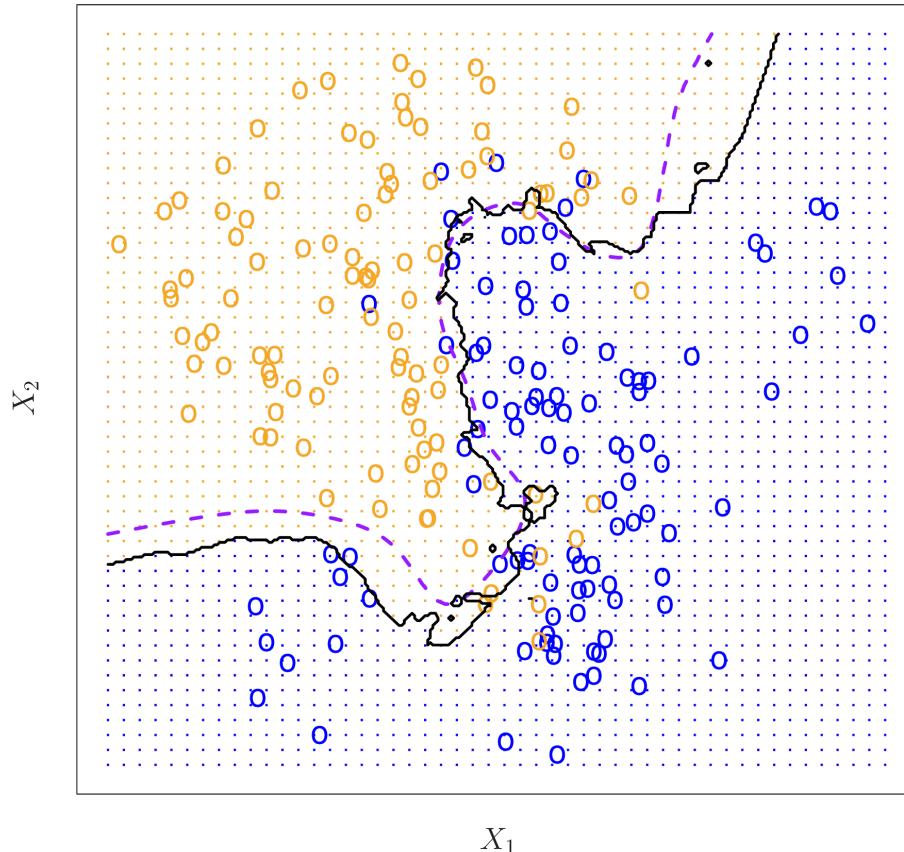
KNN: K=100



(Chapter2/2.16.pdf)

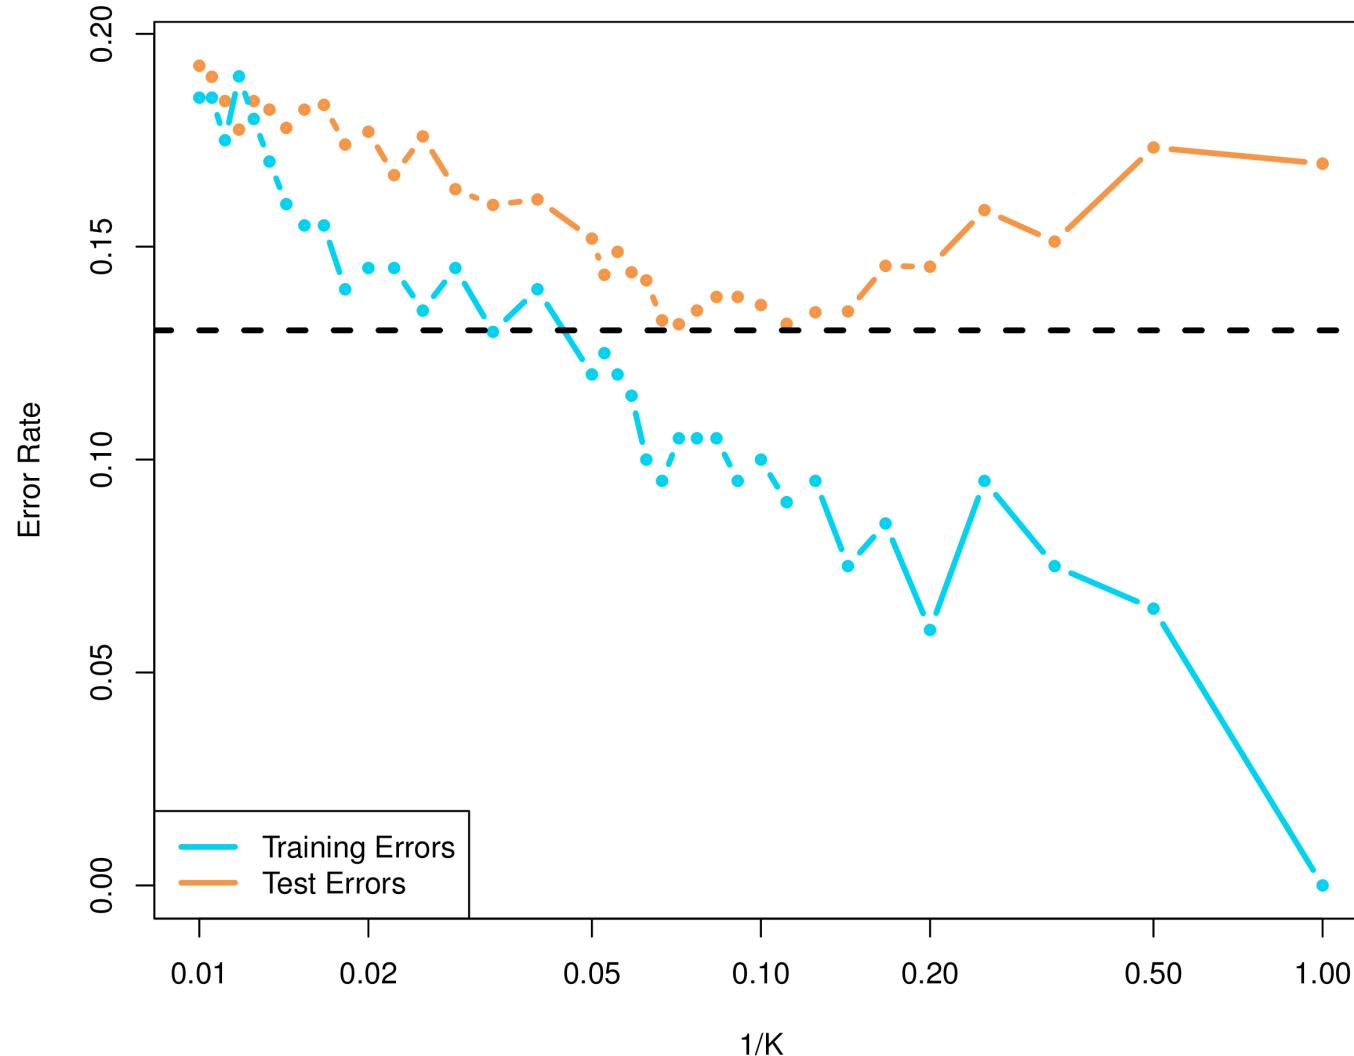
KNN: about as good as possible

KNN: K=10



(Chapter2/2.15.pdf)

KNN: bias variance trade-off, using training and test error





This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](#).

Lecturer: *Professor Di Cook*

Department of Econometrics and Business Statistics

✉ ETC3250.Clayton-x@monash.edu

CALENDAR
Week 1b

