

ETC3250/5250: Introduction to Machine Learning


Categorical response regression

Lecturer: *Professor Di Cook*

Department of Econometrics and Business Statistics

✉ ETC3250.Clayton-x@monash.edu

📅 Week 3a



Categorical responses

In **classification**, the output Y is a **categorical variable**. For example,

- ⦿ Loan approval: $Y \in \{\text{successful, unsuccessful}\}$
- ⦿ Type of business culture: $Y \in \{\text{clan, adhocracy, market, hierarchical}\}$
- ⦿ Historical document author: $Y \in \{\text{Austen, Dickens, Imitator}\}$
- ⦿ Email: $Y \in \{\text{spam, ham}\}$


Map the categories to a numeric variable, or possibly a binary matrix.

When linear regression is not appropriate


Consider the following data `simcredit` in the ISLR R package (textbook) which looks at the default status based on credit balance.



Why is a linear model not appropriate for this data?



Modelling binary responses



Orange line is a loess smooth of the data. It's much better than the linear fit.

- ⦿ To model **binary data**, we need to **link** our **predictors** to our response using a *link function*. Another way to think about it is that we will transform y to convert it to a proportion, and then build the linear model on the transformed response.
- ⦿ There are many different types of link functions we could use, but for a binary response we typically use the **logistic** link function.


The logistic function

Instead of predicting the outcome directly, we instead predict the probability of being class 1, given the (linear combination of) predictors, using the **logistic** link function.

$$p(y = 1 | \beta_0 + \beta_1 x) = f(\beta_0 + \beta_1 x)$$

where

$$f(\beta_0 + \beta_1 x) = \frac{e^{\beta_0 + \beta_1 x}}{1 + e^{\beta_0 + \beta_1 x}}$$



Logistic function

Transform the function:

$$y = \frac{e^{\beta_0 + \beta_1 x}}{1 + e^{\beta_0 + \beta_1 x}}$$

$$\longrightarrow y = \frac{1}{1/e^{\beta_0 + \beta_1 x} + 1}$$

$$\longrightarrow 1/y = 1/e^{\beta_0 + \beta_1 x} + 1$$

$$\longrightarrow 1/y - 1 = 1/e^{\beta_0 + \beta_1 x}$$

$$\longrightarrow \frac{1}{1/y - 1} = e^{\beta_0 + \beta_1 x}$$

$$\longrightarrow \frac{y}{1-y} = e^{\beta_0 + \beta_1 x}$$

$$\longrightarrow \log_e \frac{y}{1-y} = \beta_0 + \beta_1 x$$


Transforming the response makes it possible to use a linear model fit.



The left-hand side, $\log_e \frac{y}{1-y}$, is known as the **log-odds ratio** or logit. 🎲

The logistic regression model

The fitted model, where $P(Y=1|X)$ is written as: $P(Y = 1|X)$


$$\log_e \frac{P(Y=1|X)}{1-P(Y=1|X)} = \beta_0 + \beta_1 X$$

Multiple categories: This formula can be extended to more than binary response variables. Writing the equation is not simple, but follows from the above, extending it to provide probabilities for each level/category. The sum of all probabilities is 1.

Interpretation

⦿ Linear regression

- β_1 gives the average change in Y associated with a one-unit increase in X

⦿ Logistic regression

- Increasing X by one unit changes the log odds by β_1 or equivalently it multiplies the odds by e^{β_1}
- However, because the model is not linear in X , β_1 does not correspond to the change in response associated with a one-unit increase in X

Maximum Likelihood Estimation

Given the logistic $p(x_i) = \frac{1}{e^{-(\beta_0 + \beta_1 x_i)} + 1}$

We choose parameters β_0, β_1 to maximize the likelihood of the data given the model. The likelihood function is

$$l_n(\beta_0, \beta_1) = \prod_{y_i=1, i}^n p(x_i) \prod_{y_i=0, i}^n (1 - p(x_i)).$$

It is more convenient to maximize the *log-likelihood*:


$$\max_{\beta_0, \beta_1} \log l_n(\beta_0, \beta_1) = \max_{\beta_0, \beta_1} - \sum_{i=1}^n \log (1 + e^{-\beta_0 - \beta_1 x_i})$$

Making predictions

With estimates from the model fit, $\hat{\beta}_0, \hat{\beta}_1$ we can predict the **probability of belonging to class 1** using:

$$p(y = 1 | \hat{\beta}_0 + \hat{\beta}_1 x) = \frac{e^{\hat{\beta}_0 + \hat{\beta}_1 x}}{1 + e^{\hat{\beta}_0 + \hat{\beta}_1 x}}$$

Round to 0 or 1 for class prediction.



Orange points are fitted values, \hat{y} . Black points are observed response, y . (Residual is the difference between observed and predicted.) To generate categorical predictions, round the fitted values.

Fitting credit data in R

We use the `glm` function in R to fit a logistic regression model. The `glm` function can support many response types, so we specify `family="binomial"` to let R know that our response is *binary*.

```
logistic_mod <- logistic_reg() %>%  
  set_engine("glm") %>%  
  set_mode("classification") %>%  
  translate()  
  
logistic_fit <-  
  logistic_mod %>%  
  fit(default ~ balance,  
      data = simcredit)
```

Examine the fit

```
tidy(logistic_fit)
```

```
## # A tibble: 2 x 5
```

##	term	estimate	std.error	statistic	p.value
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	(Intercept)	-10.7	0.361	-29.5	3.62e-191
## 2	balance	0.00550	0.000220	25.0	1.98e-137

```
glance(logistic_fit)
```

```
## # A tibble: 1 x 8
```

##	null.deviance	df.null	logLik	AIC	BIC	deviance	df.residual	nobs
##	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>
## 1	2921.	9999	-798.	1600.	1615.	1596.	9998	10000

Write out the model

$$\hat{\beta}_0 = 10.6513306$$

$$\hat{\beta}_1 = 0.0054989$$

Model fit summary

Null model deviance 2920.6 (think of this as TSS)

Model deviance 1596.5 (think of this as RSS)

Check model fit

```
simcredit_fit <- augment(logistic_fit, simcredit)
simcredit_fit %>%
  count(default, .pred_class) %>%
  pivot_wider(names_from = "default", values_from = n)
```

```
## # A tibble: 2 x 3
##   .pred_class      No      Yes
##   <fct>          <int> <int>
## 1 No             9625    233
## 2 Yes              42    100
```




Note: Residuals not typically useful.

A warning for using GLMs!

Logistic regression model fitting fails when the data is *perfectly* separated.

MLE fit will try and fit a step-wise function to this graph, pushing coefficients sizes towards infinity and produce large standard errors.

Pay attention to warnings!



```
logistic_fit <-  
  logistic_mod %>%  
  fit(default_new ~ balance,  
      data = simcredit)
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

More on supervised classification to come

Logistic regression is a technique for supervised classification. We'll see a lot more techniques: linear discriminant analysis, trees, forests, support vector machines, neural networks.



This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/).

Lecturer: *Professor Di Cook*

Department of Econometrics and Business Statistics

✉ ETC3250.Clayton-x@monash.edu

📅 Week 3a

