MONASH University

ETC3250/5250: Introduction to Machine Learning

Classification Trees

Lecturer: *Professor Di Cook*

Department of Econometrics and Business Statistics

ETC3250.Clayton-x@monash.edu

Week 6a

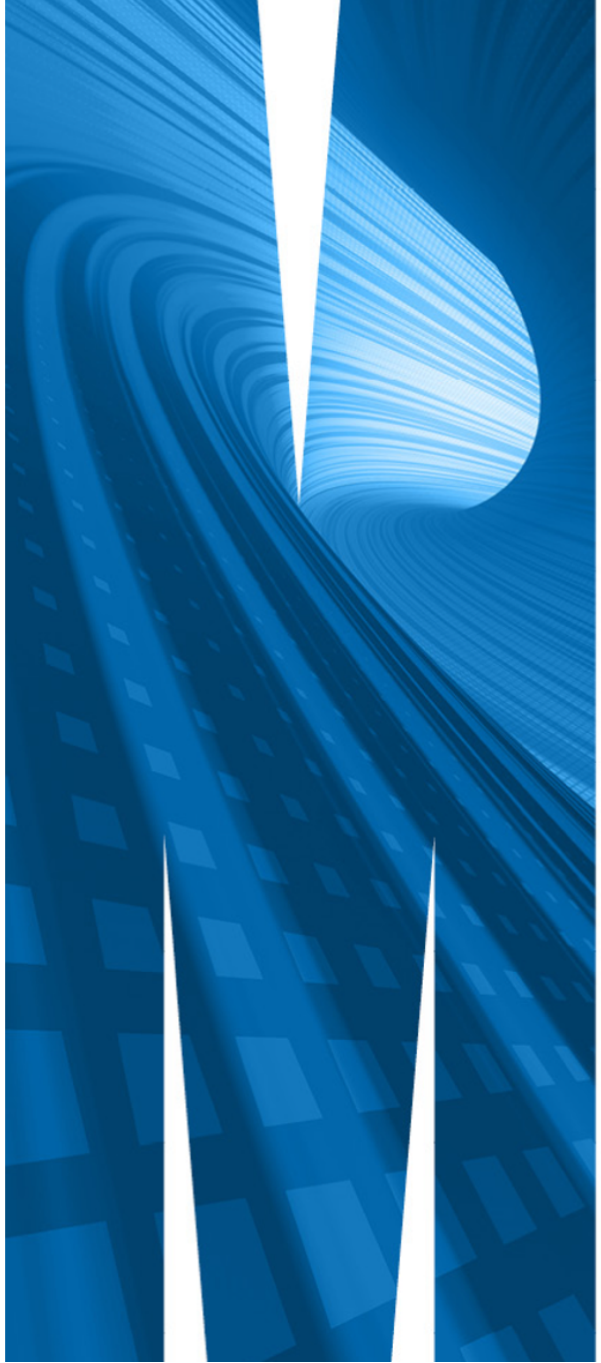

What is a decision tree?

Tree based models consist of one or more of nested `if-then` statements for the predictors that partition the data. Within these partitions, a model is used to predict the outcome.

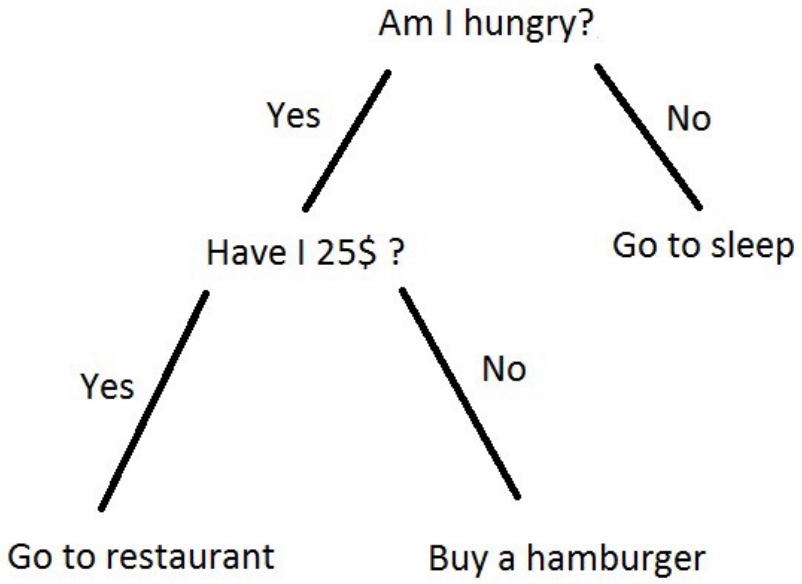

Source: Egor Dezhic

Classification trees

- A classification tree is used to predict a categorical response and regression tree is used to predict a quantitative response
- Use a recursive binary splitting to grow a classification tree. That is, sequentially break the data into two subsets, typically using a single variable each time.
- The predicted value for a new observation, x_0, will be the most commonly occurring class of training observations in the sub-region in which x_0 falls

Algorithm: growing a tree

1. All observations in a single set
2. Sort values on first variable
3. Compute split criteria for all possible splits into two sets
4. Choose the best split on this variable
5. Repeat 2-4 for all other variables
6. Choose the best split among all variables. Your data is now in two sets.
7. Repeat 1-6 on each subset.
8. Stop when stopping rule is achieved.

Split criteria - purity/impurity metrics

- The Gini index measures total variance across the K classes, for subset m:

$$G = \sum_{k=1}^{K} \hat{p}_{mk}(1 - \hat{p}_{mk})$$

- Entropy is defined as

$$D = -\sum_{k=1}^{K} \hat{p}_{mk} log(\hat{p}_{mk})$$

- If all $\hat{p}_{mk}$'s close to zero or one, G and D are small. Lower is better!

Stopping rules

- Minimum split: number of observations in a node, in order for a split to be made
- Minimum bucket: Minimum number of observations allowed in a terminal node
- Complexity parameter: minimum difference between impurity values required to continue splitting

Illustration for one variable

x	y
11	A
33	A
39	B
44	A
50	A
56	B
70	B

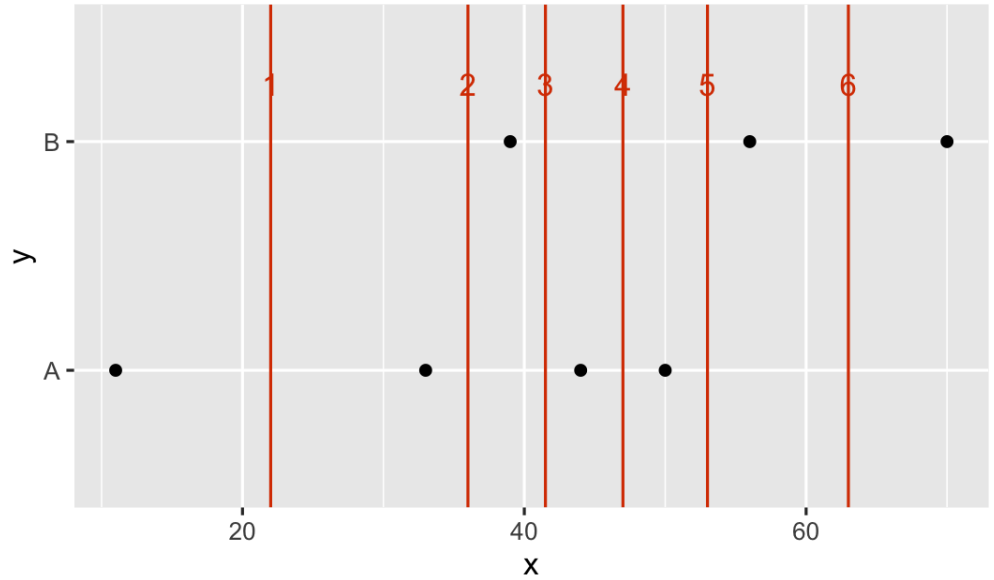

What do you think is the best split? 2, 3 or 5??

Note that x is sorted from lowest to highest!

Calculate the impurity for a split

Look at split 5.

The left bucket is

x	y
11	A
33	A
39	B
44	A
50	A

and the right bucket is

x	y
56	B
70	B

Using Gini $G = \sum_{k=1}^{K} \hat{p}_{mk}(1 - \hat{p}_{mk})$

Left bucket:

$$\hat{p}_{LA} = 4/5, \hat{p}_{LB} = 1/5, \quad p_L = 5/7$$

$$G_L = 0.8(1 - 0.8) + 0.2(1 - 0.2) = 0.32$$

Right bucket:

$$\hat{p}_{RA} = 0/2, \hat{p}_{RB} = 2/2, \quad p_R = 2/7$$

$$G_R = 0(1 - 0) + 1(1 - 1) = 0$$

Combine with weighted sum to get impurity for the split:

$$5/7 G_L + 2/7 G_R = 0.32$$

Your turn: compute the impurity for split2.

Splits on categorical variables

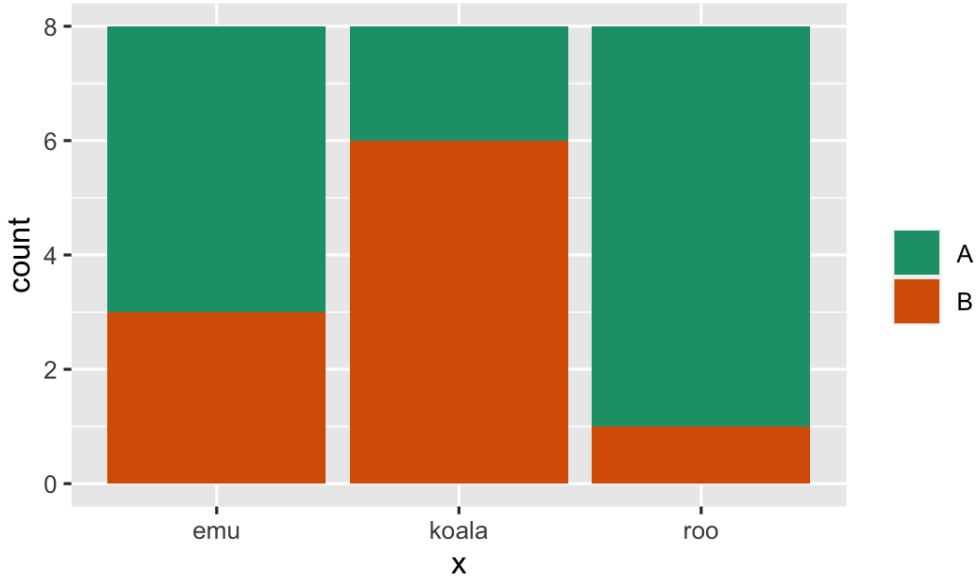

Split would be "if koala then assign to B else assign to A"

Handling missing values

x1	x2	x3	x4	y
19	-8	22	-24	A
NA	-10	26	-26	A
15	NA	32	-27	B
17	-6	27	-25	A
18	-5	NA	-23	A
13	-3	37	NA	B
12	-1	35	-30	B
11	-7	24	-31	B

50% of cases have missing values, which causes most methods to falter. For trees missings only on a single variable are ignored.

Example - predicting heart disease

Y: AHD, presence of heart disease (Yes/No)

X: heart and lung function measurements

```
##  [1] "Age"       "Sex"
##  [3] "ChestPain" "RestBP"
##  [5] "Chol"      "Fbs"
##  [7] "RestECG"   "MaxHR"
##  [9] "ExAng"     "Oldpeak"
## [11] "Slope"     "Ca"
## [13] "Thal"      "AHD"
```

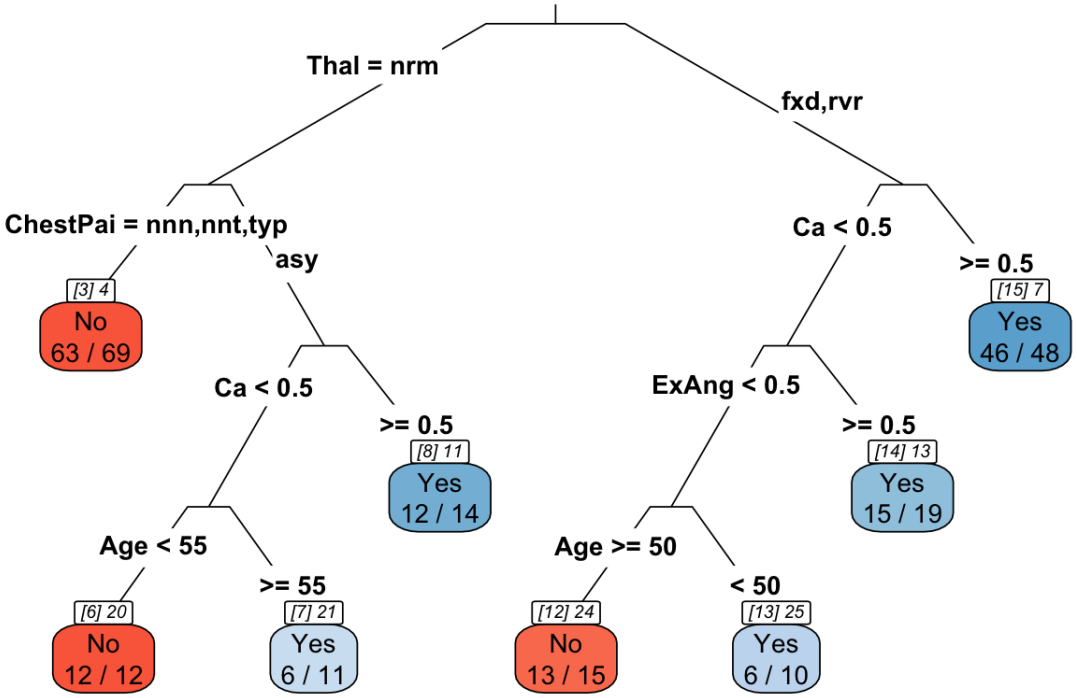

Deeper trees

Trees can be built deeper by:

- decreasing the value of the complexity parameter `cp`, which sets the difference between impurity values required to continue splitting.
- reducing the `minsplit` and `minbucket` parameters, which control the number of observations below splits are forbidden.

Larger complexity, simpler tree

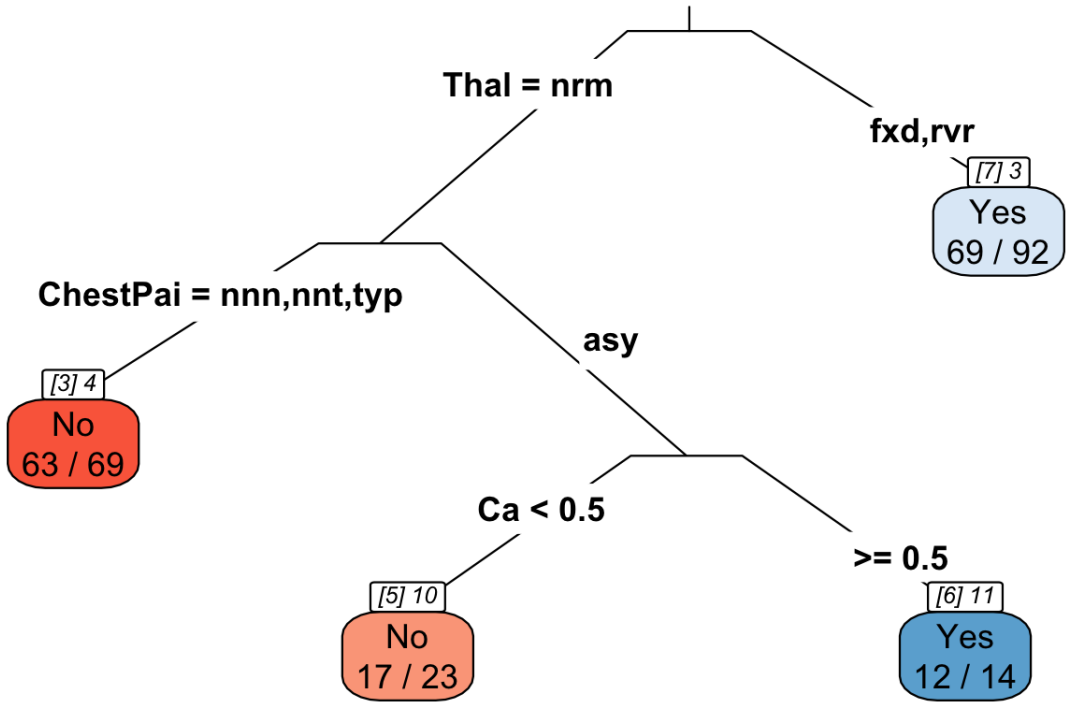

Tabulate true vs predicted to make a confusion table.

		true	
		C1 (positive)	C2 (negative)
pred-	C1	*a*	*b*
icted	C2	*c*	*d*

- Accuracy: *(a+d)/(a+b+c+d)*
- Error: *(b+c)/(a+b+c+d)*
- Sensitivity: *a/(a+c)* (true positive, recall)
- Specificity: *d/(b+d)* (true negative)
- Balanced accuracy: *(sensitivity+specificity)/2*

Training confusion and error

```
##            Truth
## Prediction No Yes
##        No  88   8
##        Yes 17  85
```

Test confusion and error

```
##            Truth
## Prediction No Yes
##        No  44   9
##        Yes 11  35
```

Training vs testing performance

- Cross-validation, 5-fold
- Grid of values in complexity, and min split

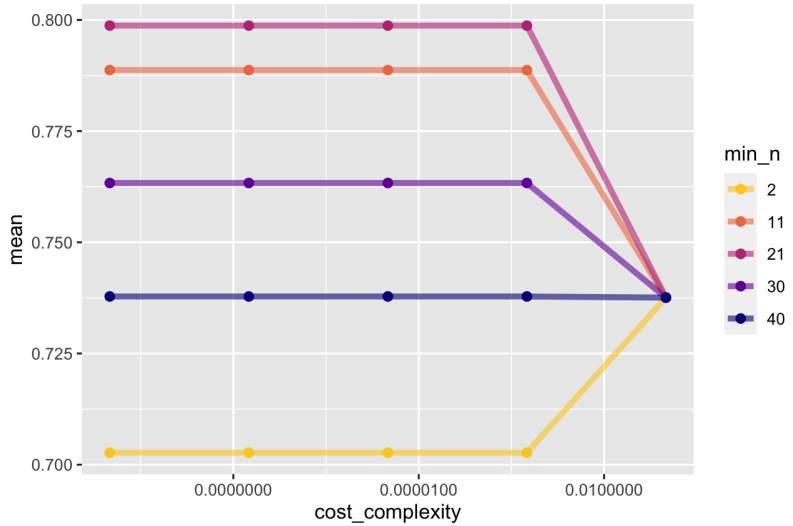


```
## # A tibble: 1 x 3
##   cost_complexity min_n
##             <dbl> <int>
## 1    0.0000000001    21
```

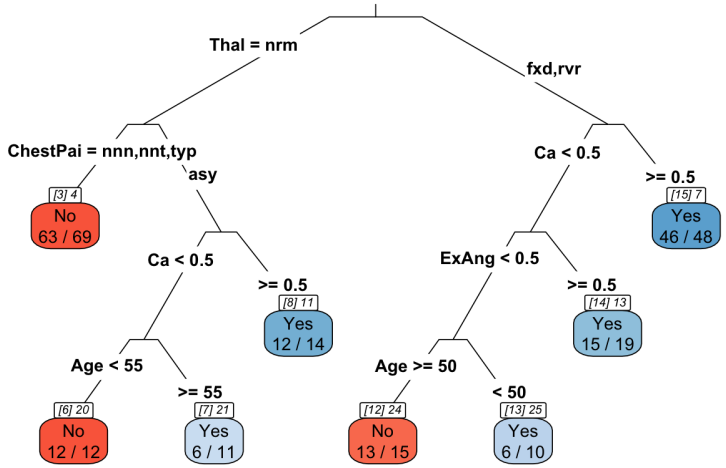

Test confusion matrix

```
##            Truth
## Prediction No Yes
##         No  44   9
##         Yes 11  35
```

Comparison with LDA

Look at the following classification problems and resultant decision boundaries for LDA (left) and CART (right).

What characteristics determine which method is more appropriate?

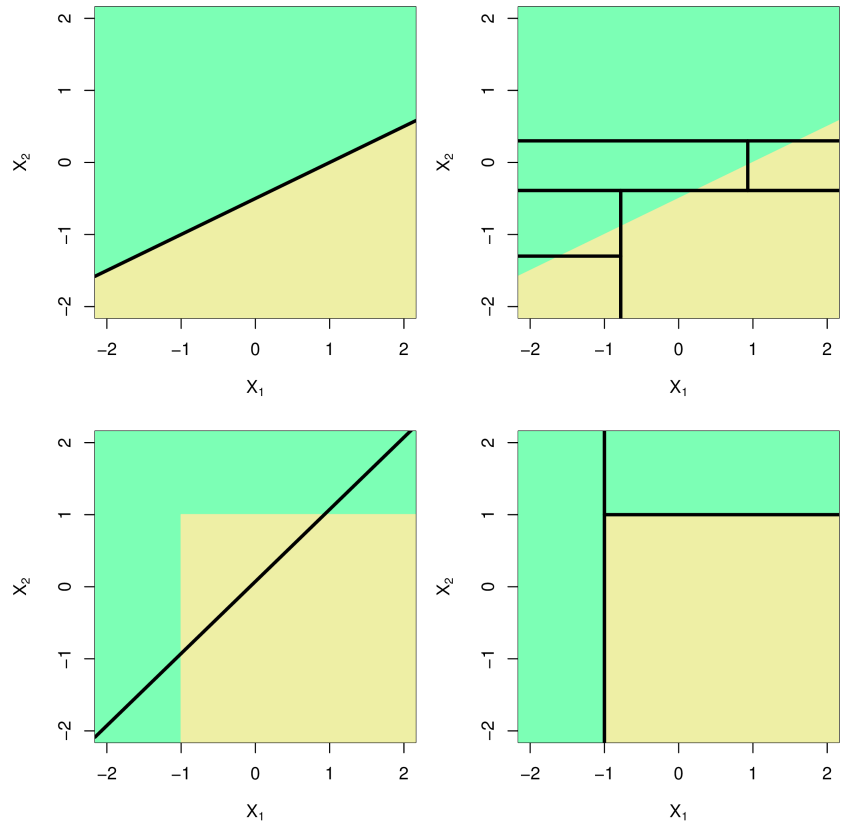

Example - Crabs

Physical measurements on WA crabs, males and females.

Data source: Campbell, N. A. & Mahon, R. J. (1974)

Decision tree parameters: minsplit=9. It's been forced to fit small subsets.

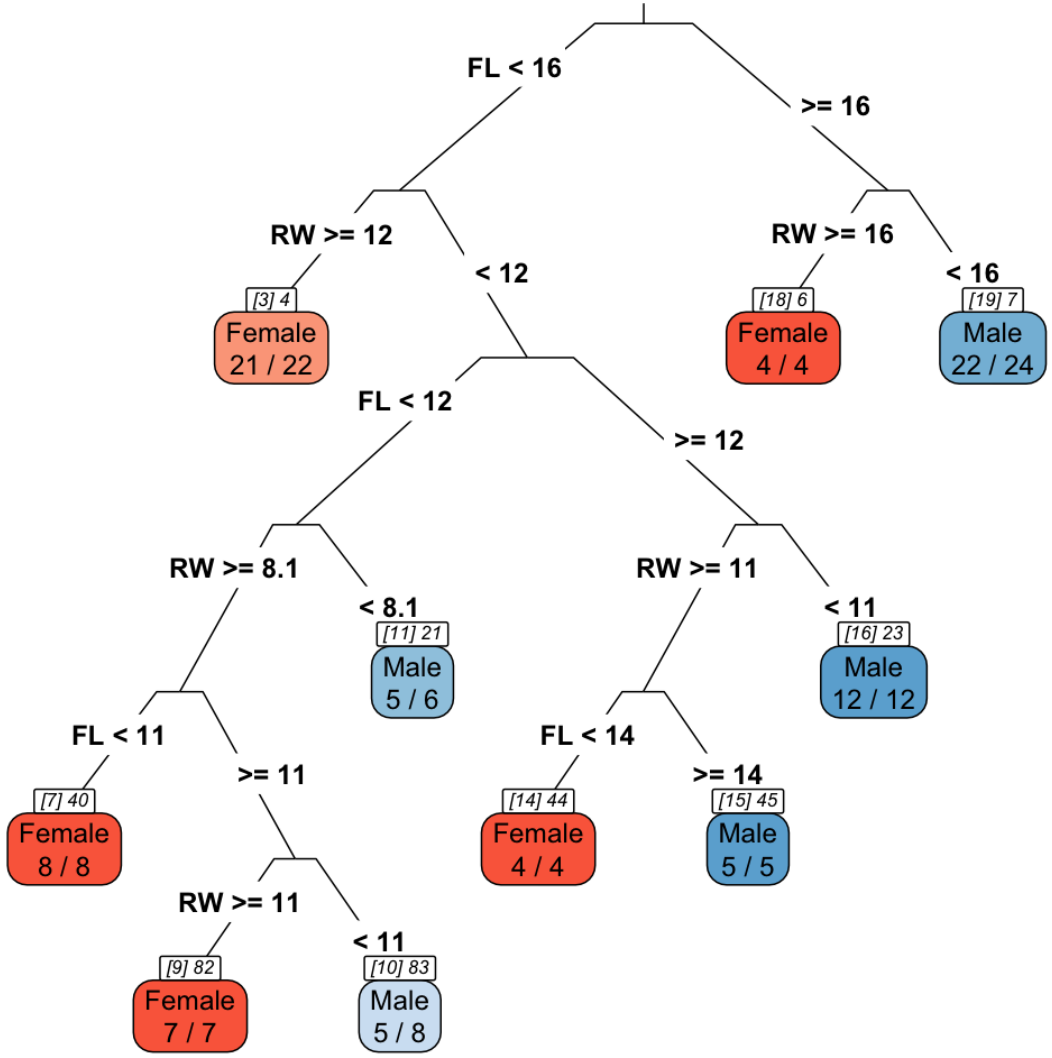

Example - Crabs

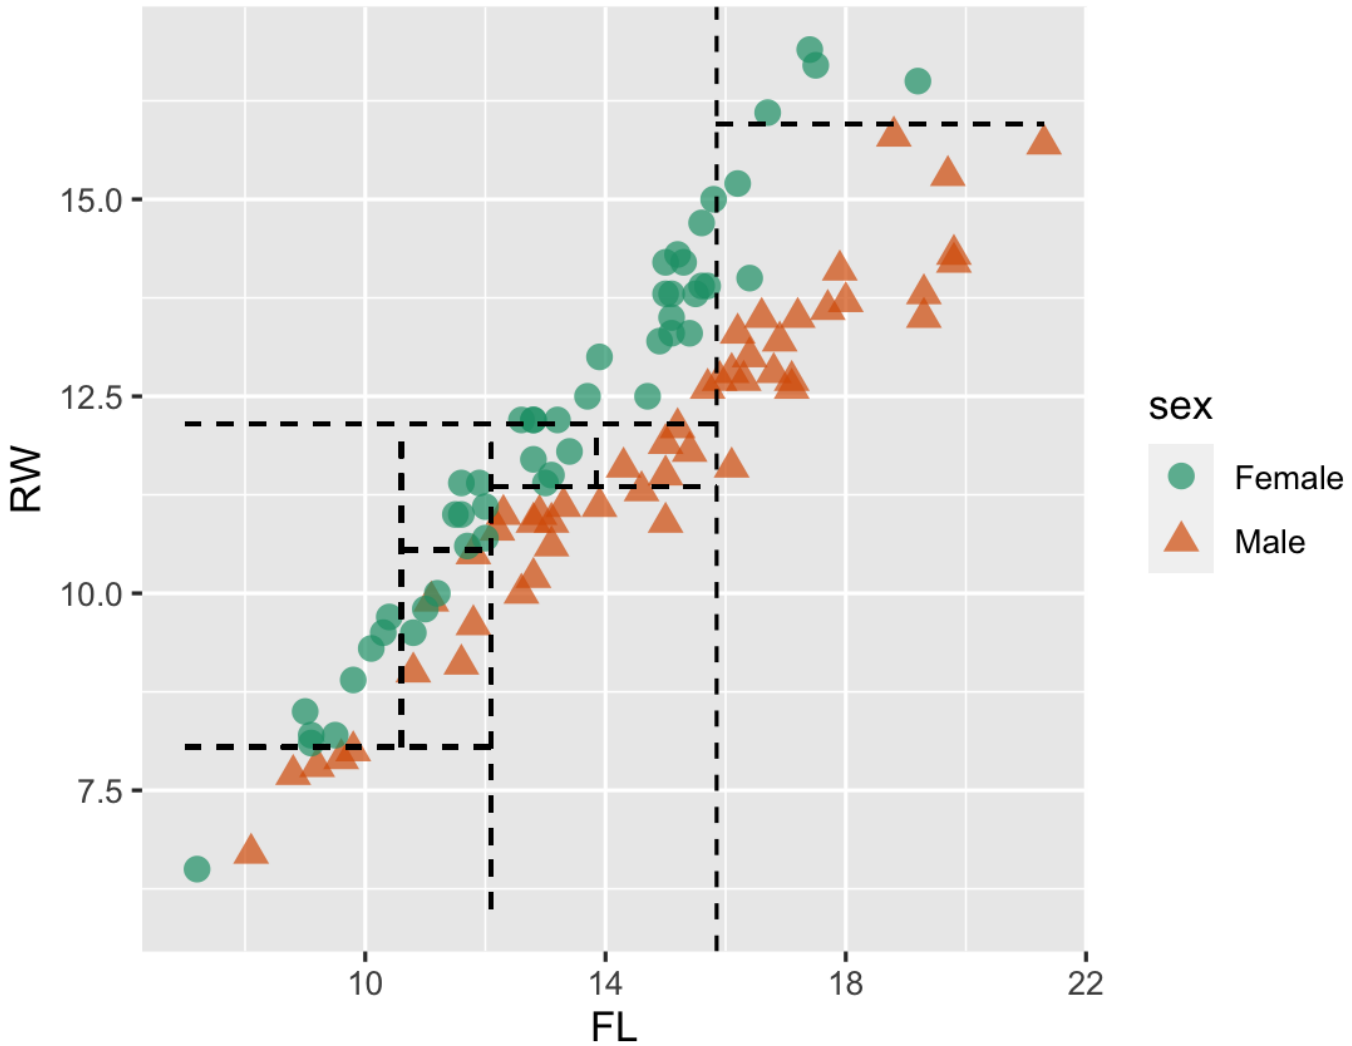

Boundaries induced by different models

layout: false

Classification tree

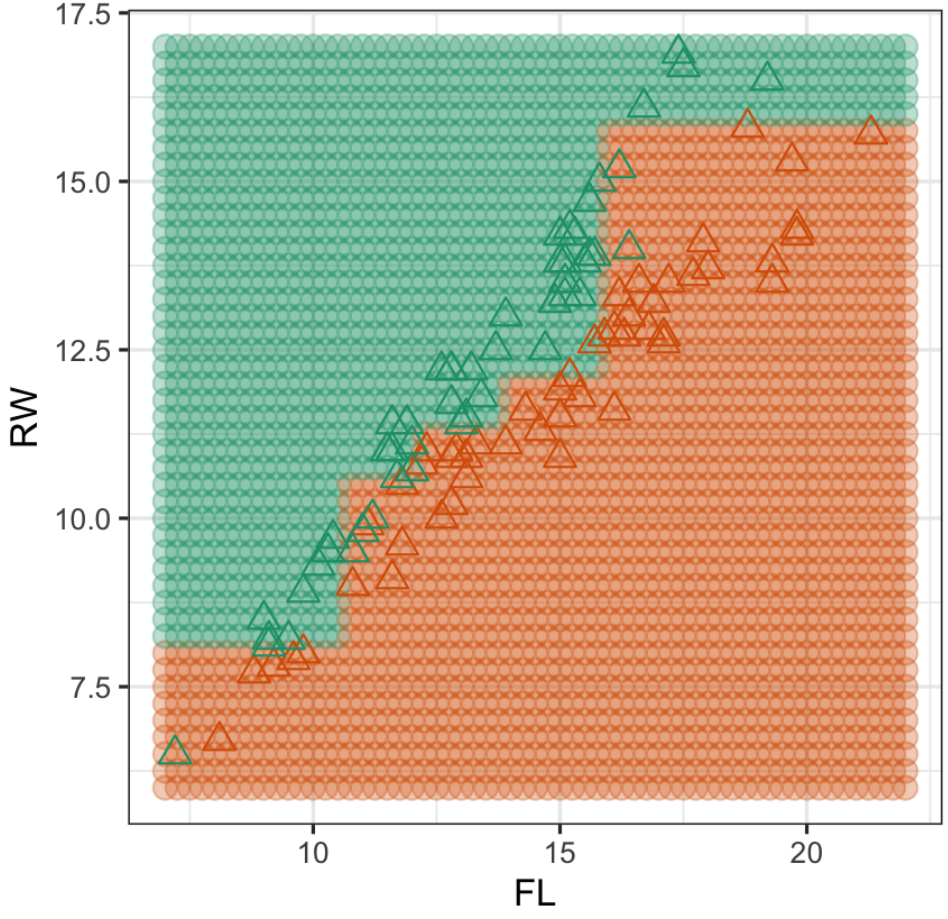

Linear discriminant classifier

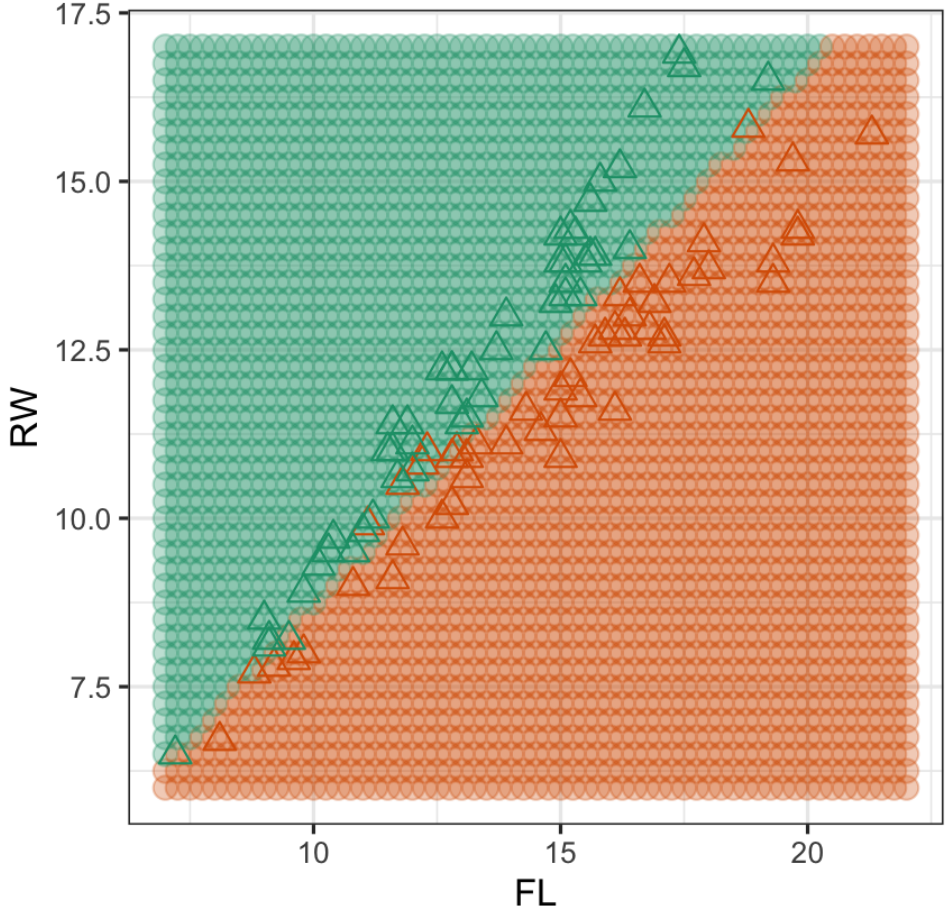

Pros and cons

- The decision rules provided by trees are very easy to explain, and follow. A simple classification model.
- Trees can handle a mix of predictor types, categorical and quantitative.
- Trees efficiently operate when there are missing values in the predictors.
- Algorithm is greedy, a better final solution might be obtained by taking a second best split earlier.
- When separation is in linear combinations of variables trees struggle to provide a good classification

MONASH University

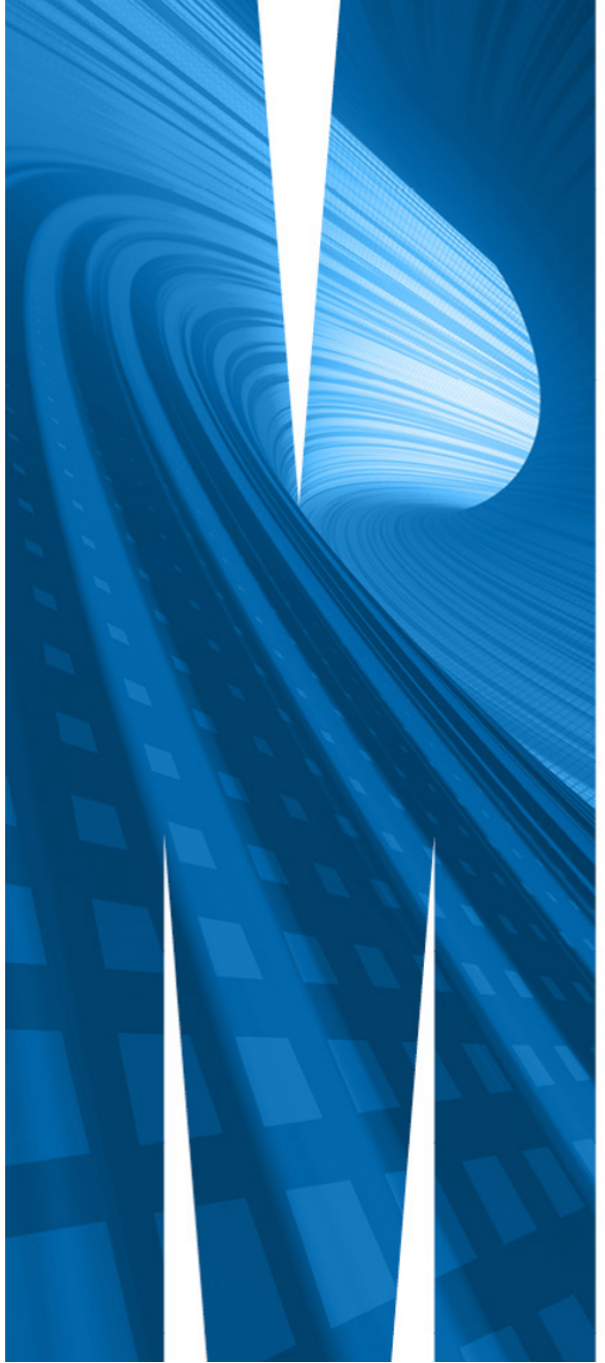

Lecturer: *Professor Di Cook*

Department of Econometrics and Business Statistics

ETC3250.Clayton-x@monash.edu

Week 6a